

FedStage DRMAA for LSF

Author: Łukasz Cieśnik lukasz.ciesnik@fedstage.com, Mariusz Mamoński mamonski@man.poznan.pl
Organization: FedStage System, Poznan Supercomputing and Networking Center
Contact: Mariusz Mamonski mamonski@man.poznan.pl
Date: 2009-05-05
Version: 1.0.5
Revision: 2340
Copyright: Copyright (C) 2007-2008 FedStage Systems, Copyright (C) 2009-2010 Poznan Supercomputing and Networking Center

Abstract

This document describes installation, configuration and usage of FedStage DRMAA for LSF version 1.0.4.

Contents

Introduction	1
Installation	1
Configuration	2
Configuration file syntax	3
Native specification	3
Release notes	5
Changes in 1.0.4 release	5
Changes in 1.0.3 release	5
Changes in 1.0.2 release	6
Changes in 1.0.1 release	6
Known bugs and limitations	7
Developers	7
Developer tools	7
Contact	7
License	8

Introduction

FedStage DRMAA for LSF is an implementation of [Open Grid Forum DRMAA 1.0](#) (Distributed Resource Management Application API) [specification](#) for submission and control of jobs to [Platform LSF](#). Using DRMAA, grid applications builders, portal developers and ISVs can use the same high-level API to link their software with different cluster/resource management systems.

This software also enables the integration of [FedStage Computing](#) with the underlying LSF system for remote multi-user job submission and control over Web Services.

Installation

To compile and install the library just go to main source directory and type:

```
$ ./configure [options] && make
$ sudo make install
```

The library requires LSF version 7.0 or further. To work with older versions it may require some patching.

Notable ./configure script options:

- with-lsf-inc LSF_INCLUDE_PATH Path to LSF header files (include dir). This and --with-lsf-lib options are unnecessary if LSF_ENVDIR environment variable is set correctly (e.g. by \$LSF_TOP/conf/profile.lsf).
- with-lsf-lib LSF_LIBRARY_PATH Path to LSF libraries (lib dir).
- with-lsf-static Link DRMAA against static LSF libraries instead of shared ones.

Note

In LSF 7.0.3 the shared libraries are broken in the way they have some undefined symbols which should be defined. Using this option fixes this problem as static libraries are built correct.

- prefix INSTALLATION_DIRECTORY Root directory where Fed-Stage DRMAA for LSF shall be installed. When not given library is installed alongside with LSF.
- enable-debug Compiles library with debugging enabled (with debugging symbols not stripped, without optimizations, and with many log messages enabled). Useful when you are to debug DRMAA enabled application or investigate problems with DRMAA library itself.

There are no unusual requirements for basic usage of library: ANSI C compiler and standard make program should suffice. If you have taken sources directly from SVN repository or wish to run test-suite you would need additional [developer tools](#). For further information regarding GNU build system see the INSTALL file.

Configuration

During DRMAA session initialization (`drmaa_init`) library tries to read its configuration parameters from locations: `/etc/lsf_drmaa.conf`, `~/lsf_drmaa.conf` and from file given in `LSF_DRMAA_CONF` environment variable (if set to non-empty string). If multiple configuration sources are present then all configurations are merged with values from user-defined files taking precedence (in following order: `$LSF_DRMAA_CONF`, `~/lsf_drmaa.conf`, `/etc/lsf_drmaa.conf`).

Currently recognized configuration parameters are:

`pool_delay` Amount of time (in seconds) between successive checks of queue(s).

Type: integer, default: 5

`cache_job_state` According to DRMAA specification every `drmaa_job_ps()` call should query DRM system for job state. With this option one may optimize communication with DRM. If set to positive integer `drmaa_job_ps()` returns remembered job state without communicating with DRM for `cache_job_state` seconds since last update. By default library conforms to specification (no caching will be performed).

Type: integer, default: 0

`wait_thread` If set to 0 every call to `drmaa_wait()` or `drmaa_synchronize()` pools DRM for selected/all jobs. By default library creates additional thread which checks state of all job for duration of DRMAA session. `drmaa_wait()/drmaa_synchronize()` calls block until finished job is found.

Type: integer, default: 1

`job_categories` Dictionary of job categories. It's keys are job categories names mapped to [native specification](#) strings. Attributes set by job category can be overridden by corresponding DRMAA attributes or native specification. Special category name default is used when `drmaa_job.category` job attribute was not set.

Type: dictionary with string values, default: empty dictionary

`lsb_events_file` The location of the `lsb.events` file. If set the library polls the LSF events logfile instead of the LSF deamons.

Type: path, default: none

Configuration file syntax

Configuration file is in form a dictionary. Dictionary is set of zero or more key-value pairs. Key is a string while value could be a string, an integer or another dictionary.

```
configuration: dictionary — dictionary_body
dictionary: '{' dictionary_body '}'
dictionary_body: (string ':' value ',')*
value: integer — string — dictionary
```

string: unquoted-string — single-quoted-string — double-quoted-string
unquoted-string: `[^ \t\n\r;0-9][^ \t\n\r;]*`
single-quoted-string: `'[^']*'`
double-quoted-string: `"[^"]*"`
integer: `[0-9]+`

Native specification

DRMAA interface allows to pass DRM dependent job submission options. Those options may be specified by settings `drmaa_native_specification` or `drmaa_job_category` job attribute. `drmaa_native_specification` accepts space delimited bsub options while `drmaa_job_category` is name of job category defined in [configuration](#) file. `-a` and `bsub` options which are meant for interactive submission of jobs (`-I`, `-Ip`, `-Is`, `-K`) are not supported.

Attributes set in native specification overrides corresponding DRMAA job attributes which overrides those set by job category.

Table 1: Native specification strings with corresponding DRMAA attributes.

DRMAA attribute	native specification
<code>drmaa_job_name</code>	<code>-J job_name</code>
<code>drmaa_input_path</code>	<code>-i input_path</code>
	<code>-is input_path</code>
<code>drmaa_output_path</code>	<code>-o output_path</code>
	<code>-oo output_path</code>
<code>drmaa_error_path</code>	<code>-e error_path</code>
	<code>-eo error_path</code>
<code>drmaa_start_time</code>	<code>-b start_time</code>
<code>drmaa_deadline_time</code>	<code>-t end_deadline</code>
<code>drmaa_js_state</code>	<code>-H</code>
<code>drmaa_transfer_files</code>	<code>-f file_stage_op</code>
<code>drmaa_v_email</code>	<code>-u mail_user</code>
	<code>-B, -N</code>
	<code>-m asked_hosts</code>
	<code>-x</code>
	<code>-n min_proc[,max_proc]</code>
	<code>-R res_req</code>
<code>drmaa_duration_hlimit</code>	<code>-c cpu_limit</code>
<code>drmaa_wct_hlimit</code>	<code>-W runtime_limit</code>
<code>drmaa_wct_slimit</code>	<code>-We estimated_runtime</code>
	<code>-M memory_limit</code>

... continued on next page

Table 1: Native specification strings with corresponding DRMAA attributes. (... continued)

DRMAA attribute	native specification
	-D data_limit
	-S stack_limit
	-v swap_limit
	-F file_limit
	-C core_limit
	-p process_limit
	-T thread_limit
	-ul
	-U reservation_id
	-ar reservation_id
	-wt warning_time
	-wa warning_action
	-s signal
	-q queue_name
	-w dependency
	-sp priority
	-r, -rn
	-G user_group
	-g job_group_name
	-P project_name
	-Lp ls_project_name
	-E pre_exec_cmd
	-Ep post_exec_cmd
	-app app_profile
	-ext sched_options
	-jsdl jsdl_doc
	-jsdl_strict jsdl_doc
	-k checkpoint_dir
	-L login_shell
	-sla service_class_name
	-Z

Release notes

Changes in 1.0.4 release

- Fixed the core limit (-C) parsing in the native specification attribute.

- Fixed infinite loop on calling `drmaa_wait/drmaa_synchronize` routines after the `CLEAN_PERIOD`

Changes in 1.0.3 release

- Fixed segfault when `drmaa_v_env` was set. Now uses `setenv` and `unsetenv` calls to modify `environ` instead of substituting `environ` pointer.
- `drmaa_transfer_files` works (in progress).
- By default when `--prefix` is not given at configure time library is installed alongside with LSF.
- When waiting for any job or with waiting thread enabled status of all jobs is pooled from DRM in one LSF API call.
- New configuration option: `cache_job_state`.
- More detailed error messages.
- It now compiles against LSF version 6.0 or further although it was not tested at runtime.

Changes in 1.0.2 release

- `drmaa_remote_command` and `drmaa_v_argv` are quoted and not interpreted by shell (e.g. spaces are allowed in command and arguments). Jobs are created with `exec` command i.e. unnecessary shell process dangling for duration of job was eliminated.
- `drmaa_wifexited` follow refinement on DRMAA Working Group mailing list - returns 1 only for exit statuses not greater than 128. Previously it returned 1 for all jobs which were run (not aborted).
- It has been reported that in some situations job which was recently submitted is not always immediately visible through LSF API. There is now workaround for such behaviour.
- `drmaa_transfer_files` is ignored because of segfaults produced by it.
- Bugfixes: Segfault when `drmaa_v_argv` is not set. Native specification parsing bugs. Various other segfaults and memory leaks.

Changes in 1.0.1 release

Note

Version 1.0.1 of library was previously released with 2.0 version number. Afterwards we decided this is misleading and does not follow versioning scheme established by DRMAA Working Group (i.e. it does not reflect the version of DRMAA specification implemented by the library).

- Many attributes implemented:

- `drmaa_start_time`,
- `drmaa_native_specification`,
- `drmaa_transfer_files`,
- job limits.
- Integrates with [FedStage Advance Reservation Library for LSF](#).
- Job category now points to native specification string in configuration file instead of job group.
- Thread safe design.
- Configuration file(s).
- Lots of bug fixes.
- More robust code.
- Meaningful logging, error messages and codes.

Known bugs and limitations

Library covers all [DRMAA 1.0 specification](#) with exceptions listed below. It was successfully tested with [Platform LSF](#) 7.0.3 on Linux OS and passes the [official DRMAA test-suite](#). All mandatory and nearly all optional job attributes (except job run duration soft limit) are implemented.

Known limitations:

- `$drmaa_incr_ph$` is replaced only within input, output and error file paths while according to specification it should be also substituted in job working directory.
- Host name is ignored in input, output and error path. They are always copied from and to submission host.
- Input file is copied from submission host when it is not present on execution host even when it was not in transfer files attribute.
- `drmaa_wcoredump()` always returns false.

Developers

Core functionality of DRMAA is put into `drmaa_utils` library. This library was created in order to keep consistent common functionality of [FedStage DRMAA for PBS Pro](#) and [FedStage DRMAA for LSF](#) library. As it is independent from any particular DRM you may find this library useful for developing other DRMAAs. For detailed information please take a look at [source code documentation](#).

Developer tools

Although not needed for library user the following tools may be required if you intend to develop FedStage DRMAA for LSF:

- GNU autotools (`autoconf`, `automake`, `libtool`),
- [Bison](#) parser generator,

- [gperf](#) perfect hash function generator,
- [Ragel](#) finite state machine compiler,
- [Docutils](#) for processing this README,
- [LaTeX](#) for creating documentation in PDF format,
- [Doxygen](#) for generating source code documentation.

Contact

Please send your comments or questions to the following mailing list:

<https://www.fedstage.com/lists/listinfo/drmaa-lsf-users> (drmaa-lsf-users@lists.fedstage.com)

Please also visit the project webpage to find news and new releases of our software:

<http://www.fedstage.com/wiki/FedStage`DRMAA`for`LSF>

License

Copyright (C) 2007-2008 FedStage Systems

Licensed under the Apache License, Version 2.0 (the “License”); you may not use this file except in compliance with the License. You may obtain a copy of the License at

<http://www.apache.org/licenses/LICENSE-2.0>

Unless required by applicable law or agreed to in writing, software distributed under the License is distributed on an “AS IS” BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied. See the License for the specific language governing permissions and limitations under the License.